




Postgres on Metal

Convex migrated from Aurora to PlanetScale Postgres Metal


AWS Aurora

PlanetScale Postgres



AWS Aurora

PlanetScale Postgres



Query p99

10–15 ms → 5–7 ms

Batch-commit p99

75–200 ms → ~20 ms

“p50 is the new p99.” Jamie Turner / Convex

Vitalize migrated from Supabase to Metal

400 GB migrated / 150 million rows

Supabase

50 s

Average read

200 GB JSONB-heavy table



PlanetScale Metal

1.2 s

Maximum read

Half the vCPUs utilized



2 ms p95



Autumn migrated to PlanetScale Metal in two days

Query latency


Previous provider

~100 ms



PlanetScale Metal

<10 ms



2 days

Migration

Seconds

Cutover downtime




Migration team

Copy + cutover support

Postgres Cluster Architecture


99.99%

single-region HA SLA




Why **Metal**: local NVMe

Network-attached storage (AWS)




PlanetScale Metal




More QPS & Lower p99 on Metal


Throughput (QPS)




p99 latency



Neki: Horizontal Scaling for Postgres



118.5M queries / second



Read-only point selects / Primary-only shards

1.22 PiB
data

512
shards

6.06 ms
router p99

Choose how you deploy PlanetScale



Google Cloud

PlanetScale-hosted

PlanetScale account



Your cloud account (BYOC)

Your AWS account / GCP project



Enterprise · PlanetScale Managed

Database operations managed by PlanetScale

Private connectivity options

Setup: PlanetScale + your team

AWS PrivateLink / GCP Private Service Connect

Migrate to PlanetScale



Self-service migration



 PlanetScale-led migration

Technical discovery + migration assessment