


Postgres on Metal






Convex migrated from Aurora to PlanetScale Postgres Metal


AWS Aurora

PlanetScale Postgres



AWS Aurora

PlanetScale Postgres



Query p99

10–15 ms → 5–7 ms

Batch-commit p99

75–200 ms → ~20 ms

"p50 is the new p99." Jamie Turner / Convex



Vitalize migrated from Supabase to Metal


400 GB migrated / 150 million rows

Supabase

50 s

Average read

200 GB JSONB-heavy table



PlanetScale Metal

1.2 s

Maximum read

Half the vCPUs utilized



2 ms p95



Autumn migrated to PlanetScale Metal in two days

Query latency

Previous provider

~100 ms



PlanetScale Metal

<10 ms



2 days

Migration

Seconds

Cutover downtime




Migration team

Copy + cutover support


Postgres Cluster Architecture

99.99%
single-region HA SLA




Why **Metal**: local NVMe

Network-attached storage (AWS)



PlanetScale Metal




More QPS & Lower p99 on Metal


Throughput (QPS)




p99 latency



Neki: Horizontal Scaling for Postgres



118.5M queries / second



1.22 PiB
data

512
shards

6.06 ms
router p99

Read-only point selects / Primary-only shards

Choose how you deploy PlanetScale



PlanetScale-hosted

PlanetScale account

 **Postgres**

Your cloud account (BYOC)

Your AWS account / GCP project

 **Postgres**

Enterprise · PlanetScale Managed

Database operations managed by PlanetScale

Private connectivity options Setup: PlanetScale + your team

AWS PrivateLink / GCP Private Service Connect

Migrate to PlanetScale

