


Postgres on **Metal**



Vitalize: 2 ms p95

Hospital staffing software

400 GB / 150 million rows

Supabase → PlanetScale Metal

PlanetScale / after migration



Half the vCPUs

vCPU

4

2

Supabase


PlanetScale

16 GB RAM on both

Postgres Cluster Architecture

99.99%

single-region HA SLA



More throughput. Lower p99

Postgres on Metal

Throughput

mean queries / second

PlanetScale 16.3k




Aurora 10.9k



AlloyDB 10.4k




Supabase 5.3k




p99 latency

milliseconds / 1-second intervals



TPCC-like / 500 GB / 32 connections / 5-minute run

Neki: Horizontal Scaling for Postgres



118.5M queries / second

[NEKI]




Read-only point selects / Primary-only shards

1.22 PiB
data

512
shards

6.06 ms
router p99

Migrate with our infrastructure team



01 Assess

Schema, extensions, workload

02 Rehearse

Replication, p95 / p99, recovery

03 Cut over

Validate, switch, monitor

Technical discovery + migration assessment


Your team + our database engineers



Convex migrated from Aurora to PlanetScale Postgres Metal


AWS Aurora

PlanetScale Postgres



AWS Aurora

PlanetScale Postgres



Query p99

10–15 ms → 5–7 ms

Batch-commit p99

75–200 ms → ~20 ms

"p50 is the new p99." Jamie Turner / Convex

Two days to PlanetScale Metal

Billing infrastructure / real-time access + credit checks



Query latency

Previous provider

~100 ms



PlanetScale Metal

<10 ms



0 25 50 75 100
milliseconds

2 days

Migration

Seconds

Cutover downtime

PlanetScale
migration team