







Convex migrated from Aurora to PlanetScale Postgres Metal


AWS Aurora

PlanetScale Postgres



AWS Aurora

PlanetScale Postgres



Query p99

10–15 ms → 5–7 ms

Batch-commit p99

75–200 ms → ~20 ms

"p50 is the new p99." Jamie Turner / Convex


Vitalize: 2 ms p95

Hospital staffing software

400 GB / 150 million rows

Supabase → PlanetScale Metal

PlanetScale / after migration



Half the vCPUs

vCPU

4

2

Supabase

PlanetScale

16 GB RAM on both



Autumn migrated to PlanetScale Metal in two days

Query latency


Previous provider

~100 ms



PlanetScale Metal

<10 ms



0

25

50

75

100

milliseconds

2 days

Migration

Seconds

Cutover downtime




Migration team

Copy + cutover support

Postgres Cluster Architecture


99.99%

single-region HA SLA




More QPS & Lower p99 on Metal


Throughput (QPS)



p99 latency




Neki: Horizontal Scaling for Postgres



118.5M queries / second

[NEKI]



Read-only point selects / Primary-only shards

1.22 PiB
data

512
shards

6.06 ms
router p99

Migrate to PlanetScale



Self-service migration



PlanetScale-led migration

Technical discovery + migration assessment