





# Vitalize: 2 ms p95

Hospital staffing software

400 GB / 150 million rows

Supabase → PlanetScale Metal

PlanetScale / after migration



Half the vCPUs

vCPU

4

2

Supabase


PlanetScale

16 GB RAM on both

# Postgres Cluster Architecture

99.99%

single-region HA SLA



# More throughput. Lower p99

Postgres on Metal

## Throughput

mean queries / second

PlanetScale 16.3k




Aurora 10.9k



AlloyDB 10.4k




Supabase 5.3k




## p99 latency

milliseconds / 1-second intervals



Mixed reads + writes / 500 GB / 32 concurrent connections / 5 minutes


# Neki: Horizontal Scaling for Postgres





# 118.5M queries / second

[ NEKI ]



Read-only point selects / Primary-only shards

1.22 PiB  
data

512  
shards


6.06 ms  
router p99



# Convex migrated from Aurora to PlanetScale Postgres Metal


AWS Aurora

PlanetScale Postgres



AWS Aurora

PlanetScale Postgres



Query p99

10–15 ms → 5–7 ms

Batch-commit p99

75–200 ms → ~20 ms

"p50 is the new p99." Jamie Turner / Convex



# Autumn migrated to PlanetScale Metal in two days

Query latency


Previous provider

~100 ms



PlanetScale Metal

<10 ms



2 days

Migration

Seconds

Cutover downtime

0

25

50

75

100

milliseconds



Migration team

Copy + cutover support



# Migrate to PlanetScale



Self-service migration



PlanetScale-led migration

Technical discovery + migration assessment