


Postgres on **Metal**



Vitalize: 2 ms p95

Hospital staffing software

400 GB / 150 million rows

Supabase → PlanetScale Metal

PlanetScale / after migration



Half the vCPUs

vCPU

4

2

Supabase


PlanetScale

16 GB RAM on both

Postgres Cluster Architecture

99.99%

single-region HA SLA



More QPS & Lower p99

Postgres on Metal

Throughput

mean queries / second

PlanetScale 16.3k




Aurora 10.9k



AlloyDB 10.4k




Supabase 5.3k




p99 latency

milliseconds / 1-second intervals



Mixed reads + writes / 500 GB / 32 concurrent connections / 5 minutes

Neki: Horizontal Scaling for Postgres



118.5M queries / second

[NEKI]



Read-only point selects / Primary-only shards

1.22 PiB
data

512
shards


6.06 ms
router p99



Convex migrated from Aurora to PlanetScale Postgres Metal


AWS Aurora

PlanetScale Postgres



AWS Aurora

PlanetScale Postgres



Query p99

10–15 ms → 5–7 ms

Batch-commit p99

75–200 ms → ~20 ms

"p50 is the new p99." Jamie Turner / Convex



Autumn migrated to PlanetScale Metal in two days

Query latency


Previous provider

~100 ms



PlanetScale Metal

<10 ms



2 days

Migration

Seconds

Cutover downtime

0

25

50

75

100

milliseconds



Migration team

Copy + cutover support

Migrate to PlanetScale



Self-service migration



PlanetScale-led migration

Technical discovery + migration assessment